

Both Sides Detected It, Neither Escalated

Concurrency and escalation failure in the July 2026 autonomous agent intrusion — a forensic reconstruction from seven public primary sources.

Detection was **not** the binding constraint at either organisation. **Escalation was.**

WHAT WE BUILT

WHAT IT PRODUCED

Ingestion pipeline over 7 cached primary sources, each parser asserting a reconciliation

16 derived datasets with content-hash provenance, 75 offline tests, and a manuscript whose every number is recomputed at build time

Escalation reconstruction from the victim's and the perpetrator's own documents

5 h 34 min from true-positive alert to response at OpenAI, with 8 further compromise events timestamped inside it

Density-restricted concurrency measure, with its analytic bound computed before the statistic

A negative result: the extent statistics were confined in advance to a band 0.016 wide. Onset ordering replaces them

EMR — an escalation-miss rate, with an ISA-18.2-style annotation that transfers between deployments

A failure mode that every current alert-screening metric scores as a *success* becomes measurable

THE HOOK

Both organisations saw it. Neither woke anyone up.

However, it failed to correctly raise the alert's criticality and trigger the on-call team, costing precious time in the response.

HUGGING FACE — ITS OWN TECHNICAL TIMELINE

The signals fired **and the correlation layer resolved them into a coherent attack signal**. Only the step from that verdict to a page failed.

OPENAI — ITS OWN 38-PAGE TECHNICAL REPORT, 19 JULY

12:03 UTC — cybersecurity monitoring raises an alert on unusual identity-related API activity, recorded in the appendix as "later correlated to ExploitGym evaluations". A true positive.

17:37 UTC — responders begin stopping the active runs.

Host root on a managed Kubernetes node at 15:59. An Artifactory administrator account created at 16:48. Both endpoints located by anchor substring on pages 37–38 of 38.

5 h 34

ALERT → RESPONSE

8

COMPROMISE EVENTS
INSIDE

2 / 2

ORGS DETECTED, NEITHER
ESCALATED

0

PRESS SOURCES THIS
RESTS ON

Detection fired. Escalation did not follow.

THE PROBLEM

The received diagnosis sends the work in the wrong direction

Every public account attributes the outcome to **detection** — the campaign was not seen soon enough. That diagnosis decides what gets built next, and the two readings diverge completely.

IF DETECTION WAS THE CONSTRAINT

More sensors. Better classifiers. Lower false-negative rates. A well-funded, well-understood programme — and none of it would have changed this outcome.

IF ESCALATION WAS THE CONSTRAINT


The sensors were already sufficient. What failed is how a **correct** verdict is ranked and routed — and nothing in the field measures that.

- **The failure is invisible to the metrics.** Alert screening is formalised over a binary label space {malicious, benign} with three failure quantities — false-positive rate, uninvestigated rate, decision-quality degradation. Each is satisfied by $\hat{y} = y$, so a correctly classified alert ranked below the paging threshold scores as a *success on every one*.
- **Process safety has what security lacks.** EEMUA 191 and ANSI/ISA-18.2 publish a target alarm-priority distribution — **80 %** low, **15 %** medium, **5 %** high — and the ASM Consortium measured **37** real operator consoles against it.
- **The checkable claim.** Security publishes no target distribution and no population study of deviation from one. Priority *is* studied; a standard against which a priority can be called *wrong* is not.

Not claimed: that Hugging Face's alarm configuration departs from 80/15/5. That configuration is unpublished. The comparison is between two research literatures, not between an organisation and a standard.

METHODOLOGY · SYSTEM ARCHITECTURE

A manuscript that cannot contradict its own artifacts



One firm finding, one negative result

PHASE	ACTIONS	ONSET	SPAN	CLASS
recon	6,191	0.014	0.986	dense
rce	2,911	0.014	0.981	dense
dropper	6,972	0.032	0.968	dense
exfil	56	0.109	0.878	sparse
c2	114	0.122	0.851	sparse
evasion	6	0.158	0.697	sparse
k8s	87	0.517	0.254	sparse
supply-chain	69	0.570	0.424	sparse
tailscale	115	0.621	0.373	sparse

ESCALATION — BOTH HALVES PRIMARY

A verbatim admission from the victim, and two anchored timestamps in OpenAI's own appendix. **Neither half depends on press reporting**, so a reader who rejects all secondary sourcing keeps this finding intact.

EXTENT VS. ONSET

Three dense phases carry **16,074** of **16,521** labelled actions and begin within **1.8 %** of campaign time of one another; the last does not begin until **62.1 %**. The record is **ordered in onset and coextensive in extent** — which defeats *stage-exclusive* response, not *stage-keyed* response.

THE NEGATIVE RESULT

Mean overlap **0.998** and coextensivity **96.3 %** were originally the headline. They are not findings: the spans alone confine them in advance to [0.984, 1] and [93.5 %, 96.8 %].

Envelopes are not duty cycles: the published summary gives only the two marginals of a phase-by-day table whose interior is not public, so simultaneous *extent* is established here and simultaneous *activity* is not.

CONCLUSION

The step from a correct verdict to a human interrupt is the one nobody measures

WHAT THE RECORD ESTABLISHES

At the victim, signals fired and were correlated into a coherent attack signal that was never raised to a criticality triggering on-call. At the perpetrator, a true-positive alert preceded the response by **5 h 34 min**. Detection was not the binding constraint at either organisation.

WHAT WE PROPOSE

EMR — an escalation-miss rate conditioned on a locally defined threshold and reported with its dual — plus an annotation of (consequence class, response window) that transfers between deployments. The adoption cost is annotation, not modelling.

WHAT WE RELEASE

The corpus, the pipeline and the verification harness, so the same measurement can be run on the next such record. Publishing the phase-by-day interior would cost no disclosure and convert our extent bound into a real duty-cycle measurement.

7

PRIMARY SOURCES, HASHED

61

CLAIMS RECOMPUTED AT BUILD

75

OFFLINE TESTS

Stated limits. One campaign, $n = 1$; no inferential statistic is reported and none would be appropriate. No value of EMR is claimed for this incident — Hugging Face's escalation policy is unpublished, so the metric is an instrument, not a measurement. Two timeline dates rest on anonymously-sourced press reporting, and every interval depending on them is flagged.

• github.com/astral-fate/agent-intrusion-escalation-forensics •

huggingface.co/datasets/FatimahEmadEldin/agent-intrusion-escalation-forensics