

APART RESEARCH × CESIA — AI INCIDENT RESPONSE SPRINT · TRACK 2

Both Sides Detected It, Neither Escalated

The July 2026 autonomous agent intrusion, reconstructed from the public record.

5 h 34

RESULT — ALERT TO
RESPONSE,
8 COMPROMISE EVENTS
INSIDE

7

DATA — PRIMARY
SOURCES,
16 DERIVED DATASETS

61

METHOD — CLAIMS
RECOMPUTED
AT BUILD TIME

0

PRESS SOURCES THE
FINDING
DEPENDS ON

• **Fatimah Emad Eldin** · AI Researcher, Trouve Labs • Full report • Deck as PDF

Both organisations detected the campaign. Why did neither wake anyone up?

It failed to correctly raise the alert's criticality and trigger the on-call team.

HUGGING FACE, ON ITS OWN DETECTION

Detection fired. Escalation did not follow.

THE PROBLEM

The diagnosis decides the remedy

THE RECEIVED DIAGNOSIS

A **detection** failure.

Remedy: more sensors, better classifiers.

WHAT THE RECORD SHOWS


An **escalation** failure.

Remedy: how a correct verdict is ranked and routed.

- A correct alert, ranked too low, scores as a **success** on every metric the field defines.
- Process safety publishes a target priority distribution — **80 / 15 / 5** — and audits **37** real consoles against it.
- Security publishes **no target**, and **no population study** of deviation from one.

SYSTEM ARCHITECTURE

Sources in, checkable claims out



METHODOLOGY

A manuscript that cannot contradict its own artifacts

PROVENANCE WITH THE DATA

Every source carries its SHA-256 and evidential class. A revision upstream shows up as a digest mismatch.

PARSERS HALT, NEVER GUESS

A partial parse looks plausible. Each stage asserts a reconciliation and stops if it fails.

NO NUMBER TYPED BY HAND

The manuscript is a template. All 61 quantities are recomputed at build; an unknown one fails the build.

- **Density restriction.** An envelope is not a duty cycle, so inference is limited to the three phases above 1,000 actions.
- **The bound comes first.** We compute what the statistic *could* have reported before reporting it.
- **Honest boundary.** This buys consistency with the sources, not correctness of the measure.

RESULTS

One firm finding, one negative result

ESCALATION — BOTH HALVES PRIMARY

12:03 alert → **17:37** response. **8** compromise events inside the gap.
Read from each organisation's own document.

ORDERED IN ONSET, COEXTENSIVE IN EXTENT

Three dense phases begin within **1.8 %** of each other; the last not until **62.1 %**. Defeats stage-*exclusive* response, not stage-keyed response.

NEGATIVE RESULT, AGAINST OURSELVES

Overlap **0.998** was the original headline. The spans alone fixed it in **[0.984, 1]** before any data was read.

PHASE	ACTIONS	ONSET	SPAN
recon	6,191	0.014	0.986
rce	2,911	0.014	0.981
dropper	6,972	0.032	0.968
exfil	56	0.109	0.878
c2	114	0.122	0.851
evasion	6	0.158	0.697
k8s	87	0.517	0.254
supply-chain	69	0.570	0.424
tailscale	115	0.621	0.373

CONCLUSION

Nobody measures the step from verdict to interrupt

WHAT THE RECORD ESTABLISHES

Detection was not the binding constraint at either organisation.

WHAT WE PROPOSE

EMR — an escalation-miss rate, with an annotation that transfers between deployments.

WHAT WE RELEASE

Corpus, pipeline and harness, so the measurement runs on the next record.

Limits. One campaign, $n = 1$. No value of EMR is claimed here — the victim's escalation policy is unpublished, so the metric is an instrument, not a measurement. Two timeline dates rest on press reporting and every interval using them is flagged.

Thank you

Questions, corrections and reanalyses are all welcome — the corpus and the harness are released so the numbers here can be checked rather than taken.

• **Fatimah Emad Eldin** · AI Researcher, Trouve Labs • Fatimah@trouve.works

[Full report](#) • [GitHub](#) • [Hugging Face dataset](#)